

Project Handoff & Developer Diary: MCPO Enhancement (Updated 2025-08-19)

Objective: To evolve the baseline `mcpo` orchestrator into a robust, manageable service with a full-featured web UI for real-time control, while aligning with the latest MCP specification and ensuring stability through testing.

Guiding Principle: Every feature, especially in the UI, must be fully functional and backed by a real backend operation. There are to be no placeholder or non-working elements.

1. Project State Snapshot (As of Handoff)

- **Overall Status:** The project is in a stable, feature-complete state based on the defined scope. All implemented UI elements are functional.
- **Test Suite:** The backend test suite is comprehensive and passing (49 passed as of this update). It covers core logic, configuration, timeout handling, structured output, health endpoint, SSL exposure, tool error envelopes, streamable HTTP transport normalization, and new persistence + read-only gating tests. Some management mutation flows remain untested (see gaps section).
- **Key Deliverables:**
 1. A hardened backend with robust error handling, timeouts, and state management.
 2. A fully integrated React-based web UI for complete operational control.
 3. Alignment with core principles of the June 2025 MCP specification.

Completed Work Breakdown:

Category	Feature	Status	Notes
Core Backend	Robust Timeout Handling	Complete	CLI flags for default/max; per-request override via header/query; returns 504 on timeout.
	MCP Spec Alignment	Complete	Injects <code>MCP-Protocol-Version</code> header; handles structured tool outputs.
	Standardized API Responses	Complete	All responses use a <code>{ok: true/false, ...}</code> envelope for predictability.
	Pydantic Config Validation	Complete	The <code>mcp.json</code> file is validated on load and reload, preventing errors from malformed configs.
	Hot Reload Safety	Complete	An <code>asyncio.Lock</code> prevents race conditions during concurrent file-watch events.
Management & API	In-Memory State Management	Complete	Tracks <code>enabled</code> state for all servers and individual tools.
	Management API (<code>/meta</code>)	Complete	Full suite of endpoints for reading state and triggering actions (reload, restart, toggle, add, remove).
	State Enforcement	Complete	Tool invocation path checks the state map and returns a 403 Forbidden if a server or tool is disabled.
	Dynamic Configuration	Complete	Servers can be added or removed via the API, which persists changes to the <code>mcp.json</code> file.
Frontend (UI)	Full UI Functionality	Complete	All toggles, buttons (Restart, Reload, Add, Remove, Open Config), and status indicators are fully wired.
	Reactivity	Complete	UI polls the backend, reflecting changes to server state and tool availability automatically.
	UI served at <code>/mcp</code>	Complete	The built React application is served directly from the Python backend.

Known Limitations / Next Steps:

1. **State Persistence (Minimal v1 Implemented):** Server/tool enabled flags persist to a versioned, atomically-written sidecar JSON state file (e.g. `<config>_state.json` / `mcpo_state.json`). On startup, stored enable maps are merged (see rules below). Only

booleans + timestamp + `version` are currently tracked (no error counts or last init metadata yet). Atomic write prevents partial/corrupt saves; unknown fields ignored for forward compatibility.

2. **"Install from Git" Feature:** The UI was designed with a tab for automated Git installation. This backend logic is **not implemented**. The UI currently only supports the "Manual Setup" path for adding new servers.
 3. **Security of Management API:** The `/_meta` endpoints that perform write actions (toggle, add, remove) are not protected by a separate, more privileged API key. If a global `--api-key` is set, it protects everything, but there is no granular control.
-

2. Developer's Diary & Rationale Log

This section details the logic behind key implementation decisions made during the project.

A. Timeout System: From Simple to Robust

- **Initial State:** The codebase had an unused `--tool-timeout` flag that did nothing.
- **Problem:** A single, global timeout is insufficient. Some tools are fast (e.g., file lookup), while others are slow (e.g., "deep research"). A low timeout would break long-running tools; a high timeout would risk hanging workers on simple calls.
- **Solution & Logic:**
 1. **Enforce the Default:** We first made the `--tool-timeout` functional by wrapping the `session.call_tool` in `asyncio.wait_for`.
 2. **Introduce a Safety Net:** We added `--tool-timeout-max` to act as a hard upper limit, preventing abuse or misconfiguration from creating excessively long timeouts.
 3. **Provide Flexibility:** The crucial step was allowing per-request overrides via the `X-Tool-Timeout` HTTP header or a `?timeout=` query parameter. This allows clients who *know* they are calling a long-running tool to request more time, up to the configured maximum.
 4. **Clear Error Handling:** A timeout now returns a distinct HTTP `504 Gateway Timeout` with a structured error `{ok: false, error: {code: "timeout"}}`, making it easy for clients to handle programmatically.

B. The "Restart" vs. "Reload" Distinction

- **Problem:** Simply re-reading the config file (`reload`) is not enough to make newly added servers functional. Starlette's `lifespan` events, which initialize the MCP connection and discover tools, only run on application startup.
- **Solution & Logic:**
 1. **Reload (`/_meta/reload`):** This is the "lightweight" action. It uses the `watchdog`-based `ConfigWatcher` to diff the config file, mounting new server sub-apps and unmounting removed ones. It's fast but doesn't initialize new servers.
 2. **Restart (`/_meta/reinit`):** This is the "heavy" action required after adding a server. We created a `reinit_servers` function that explicitly iterates through all server sub-apps and manually runs the `lifespan` startup logic for them. This connects to the MCP process, discovers tools, and makes the server fully operational without killing the main Uvicorn process. The UI "Restart" button triggers this `reinit` flow.

C. State Management & Enforcement

- **Problem:** The user required the ability to disable a misbehaving or unused server/tool instantly without removing it from the configuration.
- **Solution & Logic:**
 1. **In-Memory State:** We introduced two dictionaries on the main FastAPI `app.state`, `server_enabled` and `tool_enabled`, to act as the single source of truth. This is fast and avoids file I/O on every check.
 2. **Management API:** A suite of `.../enable` and `.../disable` endpoints were created to modify this state.
 3. **Enforcement at the Core:** The most important decision was where to enforce this state. We modified the `get_tool_handler` factory. Before any tool is called, the handler now checks the state maps. If the server or the specific tool is marked as `False`, it immediately returns an HTTP `403 Forbidden`. This is secure and efficient, as it prevents any attempt to even communicate with the underlying MCP process.
 4. **Parent App Reference:** Sub-apps needed access to the main app's state. We achieved this by passing the `parent_app` reference into each sub-app when it is mounted.

D. Frontend Philosophy: "No Placeholders"

- **Initial Conflict:** Early UI iterations included non-functional elements (like the enable/disable toggles) which led to user frustration and accusations of "glazing over" the work.

- **Corrective Action & Logic:** The development pivot was to ensure **every single interactive element in the UI was fully functional**.
 - The "Add Server" modal was initially removed because its backend didn't exist. It was then **restored** only *after* the `POST /_meta/servers` endpoint was built.
 - The toggles were re-implemented to call the real `.../enable` and `.../disable` endpoints. Their visual state is now a direct reflection of the backend's state, refreshed via polling.
 - The result is a UI with slightly fewer features than the original *design brief* (e.g., no "Install from Git"), but where **100% of what is visible to the user is real and operational**. This builds trust and provides genuine utility.
-

3. Endpoint Catalog & Capabilities

Endpoint	Method	Purpose	Mutates State	Auth (if api_key set + strict)		Idempotent	Notes
<code>/healthz</code>	GET	Basic health snapshot	No	Optional	Yes		Includes generation + last reload time
<code>/_meta/servers</code>	GET	List mounted servers + flags	No	Optional	Yes		Enabled flag derived from in-memory map
<code>/_meta/servers/{server}/tools</code>	GET	List tools for server	No	Optional	Yes		Sorted by name
<code>/_meta/config</code>	GET	Return config path	No	Optional	Yes		Path only, not contents
<code>/_meta/reload</code>	POST	Diff config + mount/unmount + init new	Yes	Optional	Partially		Existing servers untouched unless changed
<code>/_meta/reinit/{server}</code>	POST	Re-handshake a single server	Yes	Optional	No		Tears down session + rebuilds endpoints
<code>/_meta/servers</code>	POST	Add server then reload	Yes	Optional	No		Persists to config file
<code>/_meta/servers/{server}</code>	DELETE	Remove server then reload	Yes	Optional	No		Persists removal
<code>/_meta/servers/{server}/enable</code>	POST	Enable server	Yes	Optional	Yes		In-memory only currently

Endpoint	Method	Purpose	Mutates State	Auth (if api_key set + strict)	Idempotent	Notes
/_meta/servers/{server}/disable	POST	Disable server	Yes	Optional	Yes	In-memory only currently
/_meta/servers/{s}/tools/{t}/enable	POST	Enable tool	Yes	Optional	Yes	In-memory only
/_meta/servers/{s}/tools/{t}/disable	POST	Disable tool	Yes	Optional	Yes	In-memory only
{serverMount}/{tool}	POST	Invoke tool	No*	Optional	No	Underlying MCP tool call; *internal side-effects unknown

Planned (Not Implemented): metrics, version, bulk operations, state export, WebSocket/SSE updates.

4. State Model (Current vs. Proposed)

Current (persisted minimal structure in sidecar state file):

```
{
    "server_enabled": { "<server>": true/false },
    "tool_enabled": { "<server>": { "<tool>": true/false } }
}
```

Proposed enriched schema (future superset of current minimal state file):

```
{
    "version": 1,
    "servers": [
        "<server>": {
            "enabled": true,
            "tools": { "<tool>": { "enabled": true } },
            "lastInit": "2025-08-18T00:00:00Z",
            "lastError": null
        }
    ],
    "updatedAt": "2025-08-18T00:00:00Z"
}
```

Merge Rules (on startup):

1. If state file missing -> all enabled.
2. If tool missing in persisted record -> default enabled.
3. If persisted tool absent in live list -> drop silently.
4. Unknown fields ignored (forward compatibility).

5. Error & Timeout Taxonomy

Code	HTTP	Trigger	Client Guidance
<code>timeout</code>	504	Tool exceeded effective timeout	Retry with adjusted timeout if safe
<code>not_found</code>	404	Missing server or endpoint	Refresh server list; user may have removed it
<code>exists</code>	409	Add server name collision	Choose different name
<code>invalid</code>	422	Validation failure (payload/command)	Correct request; re-submit
<code>io_error</code>	500	File write failure persisting config	Check filesystem perms/disk space
<code>reload_failed</code>	500	Exception during reload apply	Inspect logs; retry once
<code>reinit_failed</code>	500	Exception during re-handshake	Retry; if persistent inspect server process
<code>no_config_mode</code>	400	Mutation when not config-driven	Run with config file or disable feature
<code>no_config</code>	400	Reload invoked without config	Provide config path
<code>disabled</code> (planned)	403	Server/tool disabled execution attempt	Enable first or skip tool

Action Item: unify current 403 responses to include code: "disabled".

6. Test Coverage Matrix (Post-Snapshot Gaps)

Domain	Scenario	Status
Enable server	Basic enable/disable + list reflects	Missing
Tool disable enforcement	Disabled tool 403	Missing
Add server (stdio)	Happy path	Missing
Add server (duplicate)	409 exists	Missing
Remove server	Tools inaccessible post-remove	Missing
Reinit	Success + preserves enable flags	Missing
Reload diff	Add/remove/update config changes	Missing
Timeout override	Below default / above max clamp	Present (base)
Structured output edge	Complex nesting classification	Partial
Concurrency	Two reload POSTs serialized	Missing
Envelope uniformity	404/422/403/500 shapes	Partial

Priority: Implement missing rows before next release. (Add tests for persistence load/merge as well.)

7. Reliability & Performance Opportunities

Short Term:

- Add exponential backoff (2,4,8...) for failed init with cap.
- Hash tool schemas to skip redundant endpoint rebuilds.
- Collapse per-tool routes into generic /invoke (optional optimization) when tool count > threshold.

Long Term:

- Replace polling with SSE/WebSocket event channel (server_added, server_removed, state_changed, tool_list_changed).
- Circuit breaker for repeatedly failing servers (auto-disable after N failures).

8. Observability Roadmap

Logging Fields (recommend): `ts`, `level`, `action`, `server`, `tool`, `duration_ms`, `outcome`, `error_code`.

Metrics (MVP): counters (`tool_calls_total`, `tool_errors_total{code}`), gauges (`servers_enabled`, `tools_enabled`), histogram (`tool_call_duration_seconds`).

Expose via: `/_meta/metrics` (JSON) or optional Prometheus format.

9. Security & Threat Model

Risks:

1. Unauthorized config mutation.
2. DoS via many long-running tool invocations (mitigated by max timeout; still concurrency risk).
3. Malicious command injection in added servers.
4. Information leakage if future verbose logs include tool inputs.

Mitigations (planned):

- Role-based API keys or signed JWT with `role=admin` for mutation endpoints.
- Command whitelist or sandbox directory for stdio servers.
- Rate limiting (token bucket) on tool execution.
- Read-only mode flag (CLI: `--read-only` or env: `MCPO_READ_ONLY=1`) (implemented) blocks all mutation endpoints.

10. Roadmap (Now / Next / Later)

Stage	Items
Now	Persistence for enable flags; mutation endpoint tests; disabled error code normalization; accountability doc
Next	Metrics endpoint; SSE updates; security (role separation); state file schema v1
Later	Circuit breaker; tool invocation history; WebSocket live logs; Git install automation; plugin discovery

11. Operational Runbook (Essential)

Reload Config: POST `/_meta/reload` → Expect `{ok:true}`; if failure inspect logs then retry once.

Add Server: POST `/_meta/servers {name, command|url,type?}` → On success tool endpoints appear after reload (auto inside handler).

Disable Misbehaving Tool: POST disable tool → Next call returns 403 (planned code `disabled`).

Recover Stuck Server: POST reinit; if still disconnected check underlying MCP process manually.

Emergency Read-Only Mode: restart with `--read-only` (or set `MCPO_READ_ONLY=1`) and all mutation endpoints return 403 (non-destructive containment mode).

12. Ethical & Professional Guidelines

1. No Placeholders: Ship only controls with real backend behavior (maintain trust).
2. Transparent Failures: Always return structured envelopes—never swallow errors silently.
3. Principle of Least Power: Default to read-only; require explicit opt-in for mutation features in production deployments.
4. Auditability: Log every mutation (who/when/what) without sensitive payload contents.
5. User Agency: Provide clear disable & rollback paths for newly added servers/tools.
6. Avoid Dark UX: No silent auto-disables—surface reasons (error count, circuit breaker) to the operator.
7. Privacy Respect: If tool inputs may contain sensitive data, allow redaction in logs (config flag).
8. Security by Configuration: Document safe defaults; insecure options must be explicit (e.g., `--allow-unsafe-commands`).
9. Minimal Attack Surface: Do not expose shell execution beyond defined MCP server commands.
10. Integrity Over Feature Velocity: Add tests before expanding surface area.

13. Immediate Low-Effort High-Impact Actions (Refreshed)

1. Add accountability mapping (UI control → endpoint) in `DECISIONS.md` or a new `ACCOUNTABILITY.md`.
2. Normalize all 403 disabled responses to ``{ok:false,error:{code:"disabled"}``.
3. Add unit tests: enable/disable, add/remove, reinit, reload diff (add/remove server), persistence merge edge (missing tool, stale server).
4. Add metrics stub endpoint returning counts (servers_enabled, tools_enabled, tool_calls_total, tool_errors_total{code}).
5. Draft design note for SSE/WebSocket event channel before implementation.

14. Handoff Verification Checklist

Item	Present	Notes
Snapshot Commit Ref	Yes	330876d7fcb200a0... (original); ongoing updates continue
Working Tests	Yes	49 passing (2025-08-19); mutation / enable-disable still needs coverage
UI Functional	Yes	All visible controls wired
Read-Only Mode Flag	Yes	--read-only / MCPO_READ_ONLY=1 blocks mutations
Mutations Protected by API Key (strict)	Optional	Single shared key; no RBAC
Error Envelope Uniformity	Improving	Disabled returns code="disabled"; audit remaining 403 paths
Persistence of Enable Flags	Partial	Minimal versioned atomic sidecar (booleans + timestamp + version)
Structured Output Classification	Basic	Images/resources tagged; circular refs -> Any
Documentation (README+Diary)	Extensive	Diary updated; consolidation later

15. Closing Note (Updated)

State persistence advanced to minimal versioned, atomic implementation and read-only mode is live. Immediate strategic priorities: (1) add tests for enable/disable + add/remove/reinit + reload diff + persistence merge edges, (2) optionally introduce a flag to disable admin endpoints entirely (in addition to read-only), (3) extend state file schema with tool/server metadata (lastInit, lastError), (4) metrics & disabled error code normalization. This diary is a living intent ledger—keep tables in sync to avoid drift.

Last editorial update: 2025-08-19.